

```
In [1]: # Load packages
import numpy as np
import pandas as pd
import statsmodels.api as sm
import matplotlib.pyplot as plt
import sklearn.linear_model as skl
from pandas import concat
from matplotlib.pyplot import subplots
from sklearn.model_selection import KFold
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LassoCV, lasso_path, RidgeCV
from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import LabelEncoder
```

```
/Users/edithsimochemo/opt/anaconda3/lib/python3.9/site-packages/panda
s/core/computation/expressions.py:21: UserWarning: Pandas requires ve
rsion '2.8.4' or newer of 'numexpr' (version '2.7.3' currently instal
led).
```

```
from pandas.core.computation.check import NUMEXPR_INSTALLED
/Users/edithsimochemo/opt/anaconda3/lib/python3.9/site-packages/panda
s/core/arrays/masked.py:60: UserWarning: Pandas requires version '1.
3.6' or newer of 'bottleneck' (version '1.3.2' currently installed).
from pandas.core import (
```

```
In [2]: ##### DATA MANIPULATION #####
```

```
In [3]: # Import dataset
data = pd.read_stata("final_fl2.dta")

# Convert Categorical variables: List of categorical columns
categorical_cols = ['race', 'sex', 'custody_description', 'county1']

# Create an instance of LabelEncoder
le = LabelEncoder()

# Apply LabelEncoder to each categorical column
for col in categorical_cols:
    data[col+ '_encoded'] = le.fit_transform(data[col])

# Select columns
data = data.select_dtypes(include=['number'])

# Removing missing observations
data = data.dropna()
data.shape
```

```
Out[3]: (129531, 90)
```

```
In [4]: # Define variables
X = data.drop(columns = ["anyrecid","distn3","distn4","finrecidany","r
Y = data["anyrecid"]

# Print the column heads
print(X.columns)

Index(['level_0', 'index', 'date', 'adate', 'rdate', 'releaseyear',
      'releasemonth', 'after', 'dist', 'distnoab', 'distn2', 'fullba
nafter',
      'fullbanbefore', 'concurrent_sentence', 'drugoffense', 'traffo
ffense',
      'otheroffense', 'smd', 'traffmar', 'traffcoc', 'traffher', 'tr
affamph',
      'traffill', 'traffconspir', 'fincrime', 'notpossoffense',
      'drugoffense_noselling', 'drugoffense_poss', 'violentcrime', '
assault',
      'elderly', 'escape', 'forgery', 'fraud', 'kidnap', 'manslaught
er',
      'murder', 'othercrime', 'otherviolent', 'propdamage', 'rackete
er',
      'robbery', 'sexcrime', 'propsteal', 'weapon', 'criminalmischie
f', 'dui',
      'licrevoke', 'fleeorescape', 'fraudforge', 'anytheft', 'anybur
g',
      'propcrime', 'birthyear', 'maxdate', 'maxadate', 'maxrdate', '
dateorig',
      'offenseyear', 'offensemonth', 'ban', 'age', 'under30', 'blac
k', 'male',
      'totalyearssentenced', 'prioroffensenum', 'prioroffense',
      'countoffenses', 'preoct97', 'placebodrug', 'placebosmd', 'yea
r_x',
      'month_x', 'unemp_rate', 'year_y', 'lincome', 'year', 'month_
y',
      'percentage_snap_recipients', 'race_encoded', 'sex_encoded',
      'custody_description_encoded', 'county1_encoded'],
      dtype='object')
```

```
In [5]: ##### OLS REGRESSION #####
```

```
In [6]: # Fit OLS
model = sm.OLS(Y,X).fit()
```

```
In [7]: # Latex Output
latex_output = model.summary().as_latex()
print(latex_output)

\begin{center}
```

```

\begin{tabular}{lclcl}
\toprule
\textbf{Dep. Variable:} & & & anyrecid & & \textbf{
R-squared:} & } & 0.201 & \backslash\backslash & & OLS & & \textbf{
\textbf{Model:} & & & & & Least Squares & & \textbf{
Adj. R-squared:} & } & 0.200 & \backslash\backslash & & & & \textbf{
\textbf{Method:} & & & & & & & \textbf{
F-statistic:} & } & 570.2 & \backslash\backslash & & & & \textbf{
\textbf{Date:} & & & & & Fri, 20 Dec 2024 & & \textbf{
Prob (F-statistic):} & } & 0.00 & \backslash\backslash & & & & \textbf{
\textbf{Time:} & & & & & 20:01:08 & & \textbf{
Log-Likelihood:} & } & -75160. & \backslash\backslash & & & & \textbf{
\textbf{No. Observations:} & & & & & 129531 & & \textbf{
AIC:} & } & 1.504e+05 & \backslash\backslash & & & & \textbf{
\textbf{Df Residuals:} & & & & & 129473 & & \textbf{
BIC:} & } & 1.510e+05 & \backslash\backslash & & & & \textbf{
\textbf{Df Model:} & & & & & 57 & & \textbf{
} & } & & \backslash\backslash & & & & \textbf{
\textbf{Covariance Type:} & & & & & nonrobust & & \textbf{
} & } & & \backslash\backslash & & & & \textbf{
\bottomrule
\end{tabular}
\begin{tabular}{lcccccc}
& & & & & \textbf{coef} & & \textbf{std}
err} & \textbf{t} & \textbf{P} > |t| & \textbf{[0.025} & \textbf{0.975]}
\midrule
\textbf{level\_0} & & & & & -3.898e-09 & & 4.53e-09
& -0.860 & & 0.390 & & -1.28e-08 & & 4.99e-09
\textbf{index} & & & & & -2.268e-09 & & 4.51e-09
& -0.503 & & 0.615 & & -1.11e-08 & & 6.57e-09
\textbf{date} & & & & & 1.0942 & & 0.205
& 5.340 & & 0.000 & & 0.693 & & 1.496
\textbf{adate} & & & & & -7.004e-05 & & 1.44e-05
& -4.872 & & 0.000 & & -9.82e-05 & & -4.19e-05
\textbf{rdate} & & & & & 0.8377 & & 0.155
& 5.397 & & 0.000 & & 0.533 & & 1.142
\textbf{releaseyear} & & & & & 0.8005 & & 0.142
& 5.631 & & 0.000 & & 0.522 & & 1.079
\textbf{releasemonth} & & & & & 0.0056 & & 0.001
& 4.152 & & 0.000 & & 0.003 & & 0.008
\textbf{after} & & & & & 0.0057 & & 0.014
& 0.412 & & 0.681 & & -0.021 & & 0.033

```

$\backslash\text{textbf{dist}}$		$\&$	$2.054e-06$	$\&$	$3.67e-05$
$\&$	0.056	$\&$	0.955	$\&$	$-6.99e-05$
$\backslash\text{textbf{distnoab}}$		$\&$	0.0235	$\&$	0.010
$\&$	2.349	$\&$	0.019	$\&$	0.004
$\backslash\text{textbf{distn2}}$		$\&$	$-8.911e-09$	$\&$	$5.39e-10$
$\&$	-16.528	$\&$	0.000	$\&$	$-9.97e-09$
$\backslash\text{textbf{fullbanafter}}$		$\&$	0.0001	$\&$	$2.43e-05$
$\&$	5.399	$\&$	0.000	$\&$	$8.37e-05$
$\backslash\text{textbf{fullbanbefore}}$		$\&$	$-4.84e-05$	$\&$	$8.96e-06$
$\&$	-5.399	$\&$	0.000	$\&$	$-6.6e-05$
$\backslash\text{textbf{concurrent_sentence}}$		$\&$	-367.3576	$\&$	68.045
$\&$	-5.399	$\&$	0.000	$\&$	-500.725
$\backslash\text{textbf{drugoffense}}$		$\&$	0.0021	$\&$	0.000
$\&$	5.392	$\&$	0.000	$\&$	0.001
$\backslash\text{textbf{traffoffense}}$		$\&$	$-1.13e-05$	$\&$	$2.09e-06$
$\&$	-5.399	$\&$	0.000	$\&$	$-1.54e-05$
$\backslash\text{textbf{otheroffense}}$		$\&$	0.0021	$\&$	0.000
$\&$	5.392	$\&$	0.000	$\&$	0.001
$\backslash\text{textbf{smd}}$		$\&$	0.0016	$\&$	0.006
$\&$	0.268	$\&$	0.788	$\&$	-0.010
$\backslash\text{textbf{traffmar}}$		$\&$	$3.645e-06$	$\&$	$6.75e-07$
$\&$	5.399	$\&$	0.000	$\&$	$2.32e-06$
$\backslash\text{textbf{traffcoc}}$		$\&$	$4.866e-06$	$\&$	$9.01e-07$
$\&$	5.399	$\&$	0.000	$\&$	$3.1e-06$
$\backslash\text{textbf{traffher}}$		$\&$	$-8.549e-06$	$\&$	$1.58e-06$
$\&$	-5.399	$\&$	0.000	$\&$	$-1.17e-05$
$\backslash\text{textbf{traffamph}}$		$\&$	$7.69e-06$	$\&$	$1.42e-06$
$\&$	5.399	$\&$	0.000	$\&$	$4.9e-06$
$\backslash\text{textbf{traffill}}$		$\&$	$-3.293e-06$	$\&$	$6.1e-07$
$\&$	-5.399	$\&$	0.000	$\&$	$-4.49e-06$
$\backslash\text{textbf{traffconspir}}$		$\&$	$-9.527e-06$	$\&$	$1.76e-06$
$\&$	-5.399	$\&$	0.000	$\&$	$-1.3e-05$
$\backslash\text{textbf{fincrime}}$		$\&$	0.0200	$\&$	0.006

&	3.346	&	0.001	&	0.008	&	0.032
\\							
\textbf{notpossoffense}				&	-0.0059	&	0.012
&	-0.509	&	0.611	&	-0.028	&	0.017
\\							
\textbf{drugoffense_noselling}				&	0.0246	&	0.013
&	1.907	&	0.056	&	-0.001	&	0.050
\\							
\textbf{drugoffense_poss}				&	0.0135	&	0.013
&	1.050	&	0.294	&	-0.012	&	0.039
\\							
\textbf{violentcrime}				&	0.0088	&	0.010
&	0.906	&	0.365	&	-0.010	&	0.028
\\							
\textbf{assault}				&	0.0020	&	0.009
&	0.219	&	0.826	&	-0.016	&	0.020
\\							
\textbf{elderly}				&	0.0198	&	0.025
&	0.806	&	0.420	&	-0.028	&	0.068
\\							
\textbf{escape}				&	0.0128	&	0.015
&	0.833	&	0.405	&	-0.017	&	0.043
\\							
\textbf{forgery}				&	0.0086	&	0.012
&	0.691	&	0.490	&	-0.016	&	0.033
\\							
\textbf{fraud}				&	-0.0190	&	0.024
&	-0.795	&	0.427	&	-0.066	&	0.028
\\							
\textbf{kidnap}				&	-0.0368	&	0.021
&	-1.764	&	0.078	&	-0.078	&	0.004
\\							
\textbf{manslaughter}				&	-0.0470	&	0.041
&	-1.153	&	0.249	&	-0.127	&	0.033
\\							
\textbf{murder}				&	-0.0680	&	0.033
&	-2.070	&	0.038	&	-0.132	&	-0.004
\\							
\textbf{othercrime}				&	0.0201	&	0.006
&	3.099	&	0.002	&	0.007	&	0.033
\\							
\textbf{otherviolent}				&	-0.0306	&	0.014
&	-2.218	&	0.027	&	-0.058	&	-0.004
\\							
\textbf{proppdamage}				&	0.0471	&	0.038
&	1.255	&	0.209	&	-0.026	&	0.121
\\							
\textbf{racketeer}				&	-0.0620	&	0.044
&	-1.409	&	0.159	&	-0.148	&	0.024
\\							

$\texttt{\textbf{robbery}}$		&	0.0151	&	0.013
& 1.126 &	0.260	&	-0.011	& 0.041	
$\texttt{\textbf{sexcrime}}$		&	0.0688	&	0.023
& 3.033 &	0.002	&	0.024	& 0.113	
$\texttt{\textbf{propsteal}}$		&	0.0667	&	0.015
& 4.383 &	0.000	&	0.037	& 0.097	
$\texttt{\textbf{weapon}}$		&	-0.0216	&	0.007
& -3.143 &	0.002	&	-0.035	& -0.008	
$\texttt{\textbf{criminalmischief}}$		&	-0.0548	&	0.041
& -1.342 &	0.180	&	-0.135	& 0.025	
$\texttt{\textbf{dui}}$		&	-0.0175	&	0.019
& -0.899 &	0.369	&	-0.056	& 0.021	
$\texttt{\textbf{licrevoke}}$		&	0.0268	&	0.008
& 3.402 &	0.001	&	0.011	& 0.042	
$\texttt{\textbf{fleeescape}}$		&	0.0373	&	0.007
& 5.048 &	0.000	&	0.023	& 0.052	
$\texttt{\textbf{fraudforge}}$		&	0.0249	&	0.012
& 2.145 &	0.032	&	0.002	& 0.048	
$\texttt{\textbf{anytheft}}$		&	0.0179	&	0.007
& 2.593 &	0.010	&	0.004	& 0.031	
$\texttt{\textbf{anyburg}}$		&	-0.0030	&	0.006
& -0.470 &	0.638	&	-0.015	& 0.009	
$\texttt{\textbf{propcrime}}$		&	-0.0297	&	0.015
& -1.929 &	0.054	&	-0.060	& 0.000	
$\texttt{\textbf{birthyear}}$		&	0.0074	&	0.004
& 1.666 &	0.096	&	-0.001	& 0.016	
$\texttt{\textbf{maxdate}}$		&	1.756e-05	&	4.56e-06
& 3.848 &	0.000	&	8.62e-06	& 2.65e-05	
$\texttt{\textbf{maxadate}}$		&	-4.127e-05	&	5.38e-06
& -7.667 &	0.000	&	-5.18e-05	& -3.07e-05	
$\texttt{\textbf{maxrdate}}$		&	-0.8381	&	0.155
& -5.400 &	0.000	&	-1.142	& -0.534	
$\texttt{\textbf{dateorig}}$		&	-1.1177	&	0.205
& -5.454 &	0.000	&	-1.519	& -0.716	

\\						
\\textbf{offenseyear}						
& 0.211 &	0.833	&	0.0105	&	0.050	
			-0.087	&		0.108
\\						
\\textbf{offensemonth}						
& 0.317 &	0.751	&	0.0013	&	0.004	
			-0.007	&		0.009
\\						
\\textbf{ban}						
& nan &	nan	&	0	&	0	
			0	&		0
\\						
\\textbf{age}						
& -0.461 &	0.645	&	-0.0020	&	0.004	
			-0.011	&		0.007
\\						
\\textbf{under30}						
& 3.187 &	0.001	&	0.0129	&	0.004	
			0.005	&		0.021
\\						
\\textbf{black}						
& 0.646 &	0.518	&	0.0330	&	0.051	
			-0.067	&		0.133
\\						
\\textbf{male}						
& 18.796 &	0.000	&	0.0348	&	0.002	
			0.031	&		0.038
\\						
\\textbf{totalyearssentenced}						
& 5.399 &	0.000	&	1.341e+05	&	2.48e+04	
			8.54e+04	&		1.83e+05
\\						
\\textbf{prioroffensenum}						
& 37.604 &	0.000	&	0.0492	&	0.001	
			0.047	&		0.052
\\						
\\textbf{prioroffense}						
& 20.341 &	0.000	&	0.0730	&	0.004	
			0.066	&		0.080
\\						
\\textbf{countoffenses}						
& 0.184 &	0.854	&	0.0008	&	0.005	
			-0.008	&		0.010
\\						
\\textbf{preoct97}						
& nan &	nan	&	0	&	0	
			0	&		0
\\						
\\textbf{placebodrug}						
& 5.392 &	0.000	&	0.0021	&	0.000	
			0.001	&		0.003
\\						
\\textbf{placebosmd}						
& 0.267 &	0.789	&	0.0016	&	0.006	
			-0.010	&		0.014
\\						
\\textbf{year_x}						
& 5.778 &	0.000	&	0.4748	&	0.082	
			0.314	&		0.636
\\						
\\textbf{month_x}						
& 2.427 &	0.015	&	0.0032	&	0.001	
			0.001	&		0.006
\\						
\\textbf{unemp_rate}						
			0.0052	&		0.001

```

& 7.368 & 0.000 & 0.004 & 0.007
\\
\textbf{year\_y} & -1.6064 & 0.304
& -5.278 & 0.000 & -2.203 & -1.010
\\
\textbf{lincome} & -0.0535 & 0.011
& -5.084 & 0.000 & -0.074 & -0.033
\\
\textbf{year} & 0.4748 & 0.082
& 5.778 & 0.000 & 0.314 & 0.636
\\
\textbf{month\_y} & 0.0032 & 0.001
& 2.427 & 0.015 & 0.001 & 0.006
\\
\textbf{percentage\_snap\_recipients} & -0.3606 & 0.037
& -9.621 & 0.000 & -0.434 & -0.287
\\
\textbf{race\_encoded} & -0.0047 & 0.017
& -0.275 & 0.783 & -0.038 & 0.029
\\
\textbf{sex\_encoded} & 0.0350 & 0.002
& 18.915 & 0.000 & 0.031 & 0.039
\\
\textbf{custody\_description\_encoded} & 0.0014 & 0.001
& 1.724 & 0.085 & -0.000 & 0.003
\\
\textbf{county1\_encoded} & 0.0002 & 6.06e-05
& 2.495 & 0.013 & 3.25e-05 & 0.000
\\
\bottomrule
\end{tabular}
\begin{tabular}{lclcl}
\textbf{Omnibus:} & & 109341.029 & \textbf{Durbin-Watson:} & \\
& 1.992 & \\
\textbf{Prob(Omnibus):} & 0.000 & \textbf{Jarque-Bera (JB):} & \\
& 9480.847 & \\
\textbf{Skew:} & 0.289 & \textbf{Prob(JB):} & \\
& 0.00 & \\
\textbf{Kurtosis:} & 1.807 & \textbf{Cond. No.} & \\
& 1.30e+16 & \\
\bottomrule
\end{tabular}
%\caption{OLS Regression Results}
\end{center}

```

Notes: \newline

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified. \newline

[2] The smallest eigenvalue is 1.92e-13. This might indicate that there are \newline

strong multicollinearity problems or that the design matrix is singular.

```
In [8]: ##### RIDGE REGRESSION #####
```

```
In [9]: # Load packages
from sklearn.linear_model import Ridge
from sklearn.metrics import mean_squared_error as mse
```

```
In [10]: # Standardized the variables
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

```
In [11]: import warnings

# Suppress all warnings
with warnings.catch_warnings():
    warnings.simplefilter("ignore")

# Set up a range of lambdas
lambdas = 10*np.linspace(5, -5, 100) / Y.std()
soln_array = skl.ElasticNet.path(X_scaled,
                                Y,
                                l1_ratio=0,
                                alphas=lambdas)[1]

print(soln_array.shape)
```

(84, 100)


```
In [12]: # Transform Soln_path into a dataframe
soln_path = pd.DataFrame(soln_array.T,
                          columns=X.columns,
                          index=-np.log(lambdas))
soln_path.index.name = 'negative log(lambda)'
soln_path
```

Out[12]:

	ate	year_y	lincome	year	month_y	percentage_snap_recipients	race_encoded
	9e-07	-8.073586e-07	-3.510545e-07	-8.073586e-07	1.855850e-09	-7.588481e-07	-3.429147e-07
	8e-07	-1.018752e-06	-4.429712e-07	-1.018752e-06	2.341490e-09	-9.575399e-07	-4.327054e-07
	9e-07	-1.285488e-06	-5.589514e-07	-1.285488e-06	2.954104e-09	-1.208250e-06	-5.460063e-07
	5e-07	-1.622052e-06	-7.052930e-07	-1.622052e-06	3.726828e-09	-1.524594e-06	-6.889722e-07
	9e-06	-2.046719e-06	-8.899411e-07	-2.046719e-06	4.701403e-09	-1.923747e-06	-8.693690e-07

	1e-02	8.655230e-03	-7.427179e-03	-9.063092e-04	1.257519e-03	-4.118610e-02	-7.989319e-03
	0e-02	9.832555e-03	-7.428573e-03	-4.928268e-04	1.397270e-03	-4.117963e-02	-7.953269e-03
	0e-02	1.089884e-02	-7.429723e-03	-1.407429e-04	1.524582e-03	-4.117455e-02	-7.921271e-03
	5e-02	1.185477e-02	-7.430680e-03	1.562518e-04	1.639241e-03	-4.117054e-02	-7.892679e-03
	0e-02	1.270382e-02	-7.431485e-03	4.046866e-04	1.741431e-03	-4.116730e-02	-7.866986e-03

```
In [13]: # Plot the graph
path_fig, ax = subplots(figsize=(8,8))
soln_path.plot(ax=ax, legend=False)
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Standardized coefficients', fontsize=20)
#ax.legend(loc='upper left', bbox_to_anchor=(1.1, 1.005))
#ax.set_title('Ridge Regression Coefficient Path', fontsize=21);
plt.savefig('ridge_final.png', format = 'png', dpi=300, bbox_inches='t')
#ax.set_ylim([-0.1,0.4]);
```



```
In [14]: import matplotlib.pyplot as plt
from matplotlib.colors import to_rgba

# Plot the graph
fig, ax = plt.subplots(figsize=(8, 8))

# Plot all the coefficient paths
soln_path.plot(ax=ax, legend=False)
```

```

# List of variables you want to highlight with their specific colors and
highlight_vars = {
    'income': ('black', 'Income'),
    'unemp_rate': ('blue', 'Unemployment'),
    'dist': ('green', 'Days from Cutoff'),
    'percentage_snap_recipients': ('red', 'SNAP Recipients')
}

# Iterate over the plotted lines
for i, line in enumerate(ax.get_lines()):
    var_name = soln_path.columns[i]

    # Check if the current variable is one to highlight
    if var_name in highlight_vars:
        color, label = highlight_vars[var_name]
        line.set_color(color)
        line.set_linewidth(2)


        # Annotate the line
        x_data = line.get_xdata()[-1]
        y_data = line.get_ydata()[-1]
        ax.text(x_data, y_data, label, color=color, fontsize=12, fontweight='bold',
                verticalalignment='center', horizontalalignment='left')
    else:
        # Fade the other variables
        line.set_color(to_rgba('gray', alpha=0.4))
        line.set_linewidth(1)


# Set labels and title
ax.set_xlabel('$-\log(\lambda)$', fontsize=15)
ax.set_ylabel('Standardized coefficients', fontsize=15)
#ax.set_title('Ridge Regression Coefficient Path', fontsize=21)
ax.legend().set_visible(False)

# save the figure
plt.savefig('ridge_faded.png', format='png', dpi=300, bbox_inches='tight')

plt.show();

```





```
In [15]: # Ridge cross-validation plot
# Set up cross validation
K = 5
kfold = KFold(n_splits = K, random_state=0, shuffle=True)
```

In [16]: **import** warnings

Suppress all warnings

with warnings.catch_warnings():
 warnings.simplefilter("ignore")

Perform RidgeCV with different alpha (lambda) values


ridgeCV = skl.ElasticNetCV(alphas=lambdas,
 l1_ratio=0,
 cv=kfold)

pipeCV = Pipeline(steps=[('scaler', scaler),
 ('ridge', ridgeCV)])

print(pipeCV.fit(X, Y))

```
Pipeline(steps=[('scaler', StandardScaler()),
                 ('ridge',
                  ElasticNetCV(alphas=array([2.06826854e+05, 1.6390674
5e+05, 1.29893292e+05, 1.02938213e+05,
8.15767731e+04, 6.46481976e+04, 5.12325910e+04, 4.06009522e+0
4,
3.21755603e+04, 2.54985813e+04, 2.02071896e+04, 1.60138522e+0
4,
1.26907040e+04, 1.00571659e+04, 7.97013196e+03, 6.31619328e+0
3,
5.00547516e+03, 3.96675346e+03,...
1.71711288e-03, 1.36078259e-03, 1.07839693e-03, 8.54611126e-0
4,
6.77264702e-04, 5.36720694e-04, 4.25341971e-04, 3.37076238e-0
4,
2.67127154e-04, 2.11693701e-04, 1.67763638e-04, 1.32949814e-0
4,
1.05360454e-04, 8.34963580e-05, 6.61694358e-05, 5.24381462e-0
5,
4.15563341e-05, 3.29326841e-05, 2.60985890e-05, 2.06826854e-0
5])),
                 cv=KFold(n_splits=5, random_state=0, sh
uffle=True),
                 l1_ratio=0))])
```

```
In [17]: tuned_ridge = pipeCV.named_steps['ridge']
ridgeCV_fig, ax = subplots(figsize=(8,8))
ax.errorbar(-np.log(lambdas),
            tuned_ridge.mse_path_.mean(1),
            yerr=tuned_ridge.mse_path_.std(1) / np.sqrt(K))
ax.axvline(-np.log(tuned_ridge.alpha_), c='k', ls='--')
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Cross-validated MSE', fontsize=20);
```



```
In [18]: ##### LASSO REGRESSION #####
```

```
In [19]: # Lasso
import warnings
import matplotlib.pyplot as plt

# Suppress all warnings
with warnings.catch_warnings():
    warnings.simplefilter("ignore")

    lassoCV = skl.ElasticNetCV(n_alphas=200, l1_ratio=1, cv=kfold)
    pipeCV = Pipeline(steps=[('scaler', scaler), ('lasso', lassoCV)])
    pipeCV.fit(X, Y)
    tuned_lasso = pipeCV.named_steps['lasso']
    best_alpha = tuned_lasso.alpha_
```

```
In [20]: # Compute Lasso path
# lambdas, soln_array = lasso_path(X_scaled, Y.values.ravel(), alphas=r

lambdas, soln_array = lasso_path(X_scaled, Y,
                                alphas=np.logspace(-4, 4, 100),
                                max_iter=10000)[:2]
soln_path = pd.DataFrame(soln_array.T,
                        columns=X.columns,
                        index=-np.log(lambdas))

soln_path
```

Out[20]:

ab	...	unemp_rate	year_y	lincome	year	month_y	percentage_snap_recipients	race_encoded
30	...	-0.000000	-0.0	-0.000000	-0.0	0.0	-0.000000	-0.000000
30	...	-0.000000	-0.0	-0.000000	-0.0	0.0	-0.000000	-0.000000
30	...	-0.000000	-0.0	-0.000000	-0.0	0.0	-0.000000	-0.000000
30	...	-0.000000	-0.0	-0.000000	-0.0	0.0	-0.000000	-0.000000
30	...	-0.000000	-0.0	-0.000000	-0.0	0.0	-0.000000	-0.000000
...
37	...	0.015556	-0.0	-0.007064	-0.0	0.0	-0.043357	-0.006769
73	...	0.015444	-0.0	-0.007121	-0.0	0.0	-0.042943	-0.006841
14	...	0.015352	-0.0	-0.007168	-0.0	0.0	-0.042604	-0.006911
42	...	0.015276	-0.0	-0.007207	-0.0	0.0	-0.042324	-0.007021
17	...	0.015210	-0.0	-0.007239	-0.0	0.0	-0.042085	-0.007126

```
In [21]: # Plot the solution path
```



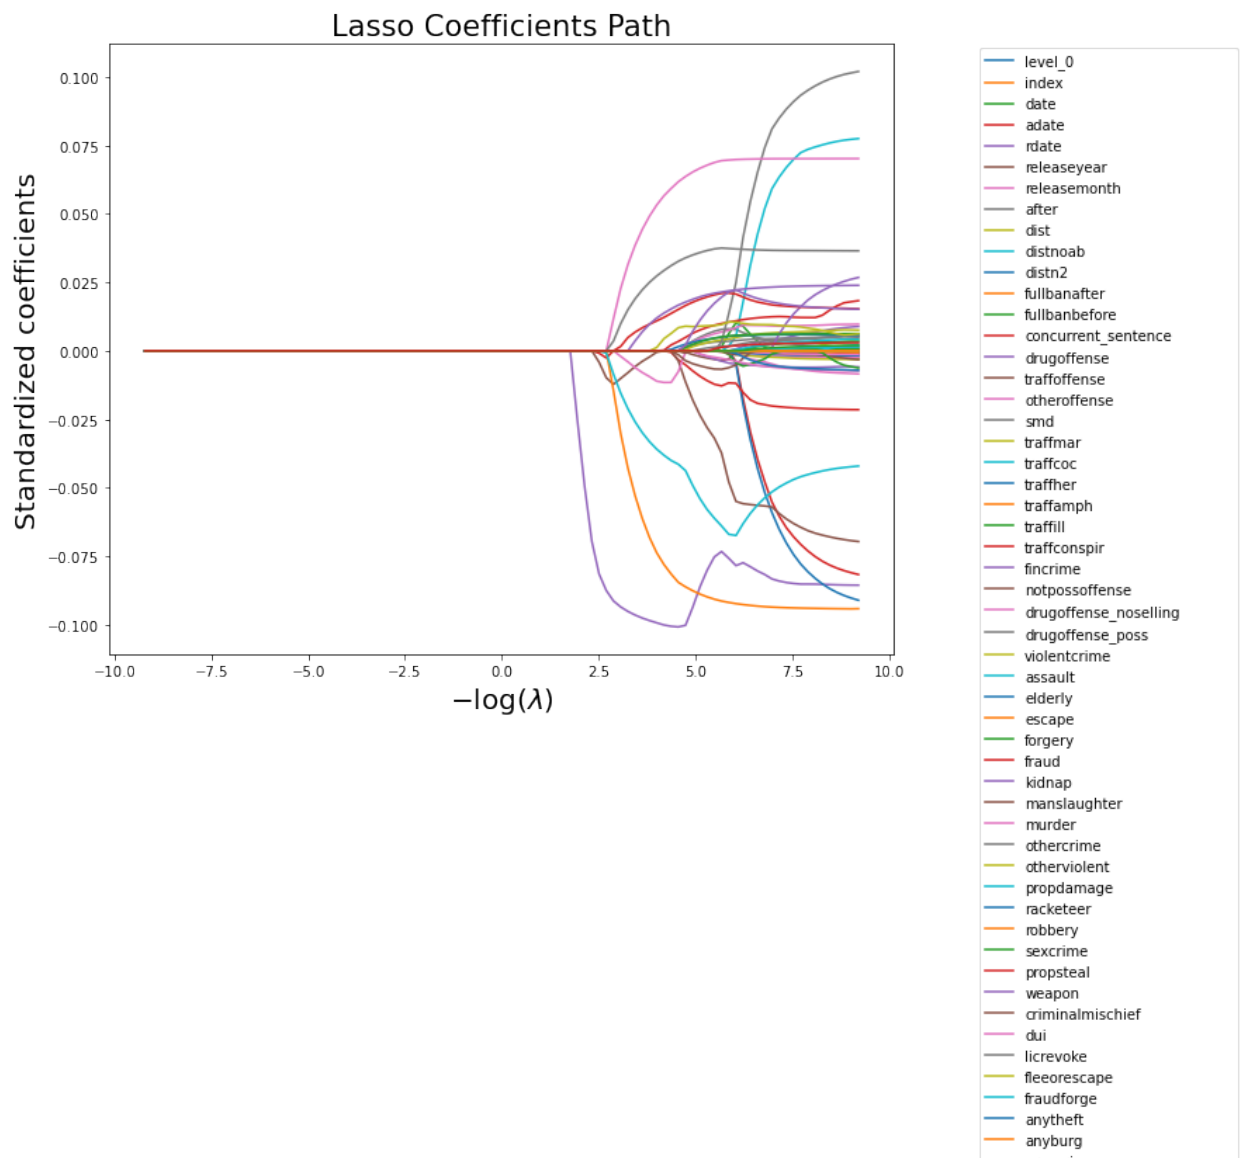
```
plt.figure(figsize=(10, 8))
ax = plt.gca()

for column in soln_path.columns:
    ax.plot(soln_path.index.to_numpy(), soln_path[column].to_numpy(),

ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Standardized coefficients', fontsize=20)
ax.set_title('Lasso Coefficients Path', fontsize=21)
ax.legend(loc='upper left', bbox_to_anchor=(1.1, 1.005))
plt.tight_layout()
plt.show();
```

/var/folders/1x/fvyxnz7d5db2mw8jz8k69snr0000gn/T/ipykernel_48306/2188380166.py:12: UserWarning: Tight layout not applied. The bottom and top margins cannot be made large enough to accommodate all axes decorations.

```
plt.tight_layout()
```



propcrime
birthyear
maxdate
maxadate
maxrdate
dateorig
offenseyear
offensemonth
ban
age
under30
black
male
totalyearssentenced
prioroffensenum
prioroffense
countoffenses
preoct97
placebodrug
placebosmd
year_x
month_x
unemp_rate
year_y
lincome
year
month_y
percentage_snap_recipients
race_encoded
sex_encoded
custody_description_encoded
county1_encoded

```
In [22]: # Cross validation MSE
import warnings

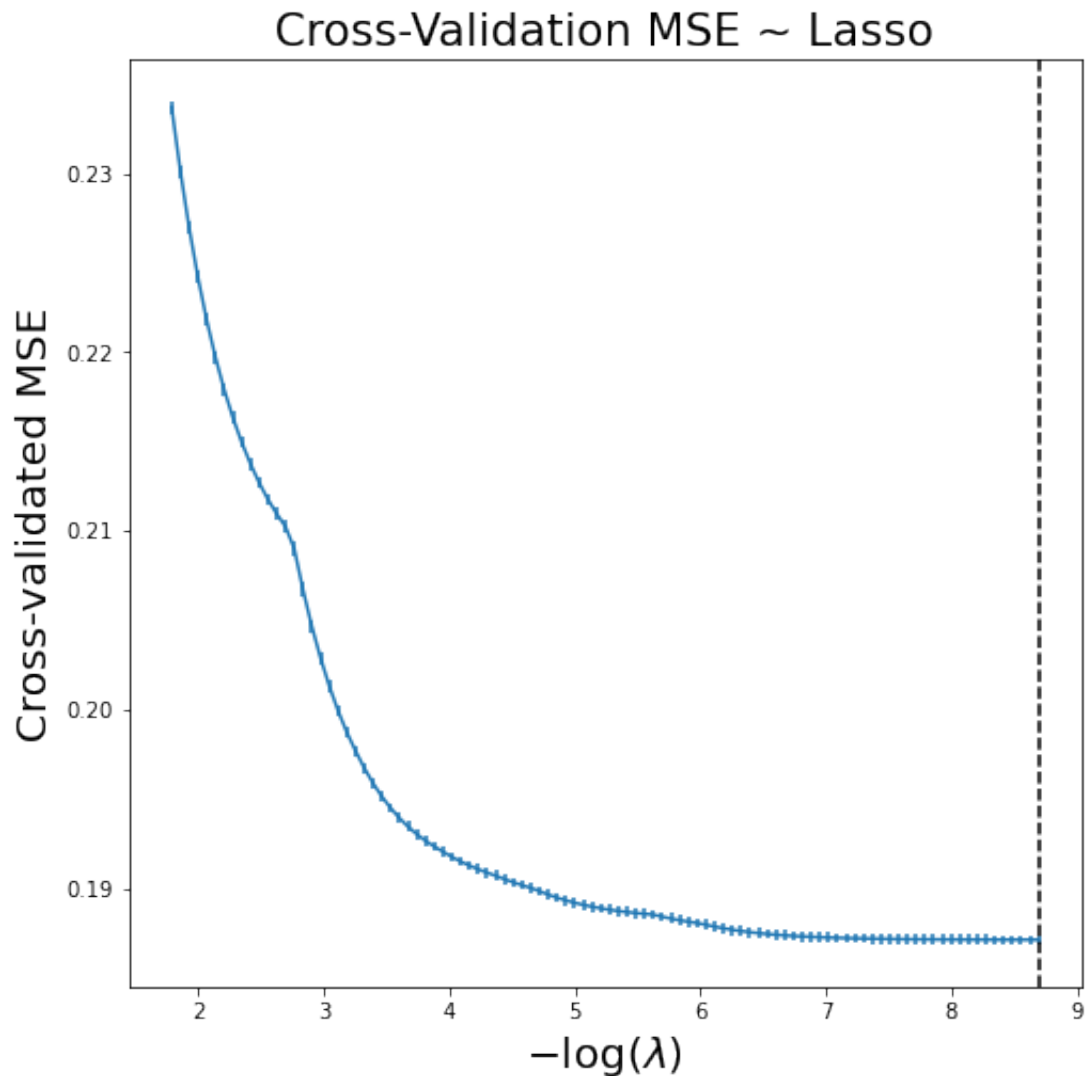
# Suppress all warnings
with warnings.catch_warnings():
    warnings.simplefilter("ignore")

    lassoCV = skl.ElasticNetCV(n_alphas=100,
                              l1_ratio=1,
                              cv=kfold)

    pipeCV = Pipeline(steps=[('scaler', scaler),
                              ('lasso', lassoCV)])

    pipeCV.fit(X, Y)
    tuned_lasso = pipeCV.named_steps['lasso']
    tuned_lasso.alpha_
```

```
In [23]: lassoCV_fig, ax = subplots(figsize=(8,8))
ax.errorbar(-np.log(tuned_lasso.alphas_),
            tuned_lasso.mse_path_.mean(1),
            yerr=tuned_lasso.mse_path_.std(1) / np.sqrt(K))
ax.axvline(-np.log(tuned_lasso.alpha_), c='k', ls='--')
ax.set_xlabel('$-\log(\lambda)$', fontsize=20)
ax.set_ylabel('Cross-validated MSE', fontsize=20)
ax.set_title('Cross-Validation MSE ~ Lasso', fontsize =21);
```



```
In [24]: ##### CLASSIFICATION TREE ###
```

```
In [25]: # Load packages
import sklearn.model_selection as skm
from sklearn.tree import (DecisionTreeClassifier as DTC,
                          DecisionTreeRegressor as DTR,
                          plot_tree,
                          export_text)
from sklearn.metrics import (accuracy_score,
                             log_loss)
from sklearn.ensemble import \
    (RandomForestRegressor as RF,
     GradientBoostingRegressor as GBR,
     GradientBoostingClassifier as GBC,
     RandomForestClassifier as RFC)
```

```
In [26]: # classification Tree
clf = DTC(criterion='entropy',
          max_depth=3,
          random_state=0)
clf.fit(X,Y)
```

```
Out[26]: DecisionTreeClassifier
DecisionTreeClassifier(criterion='entropy', max_depth=3, random_state=0)
```

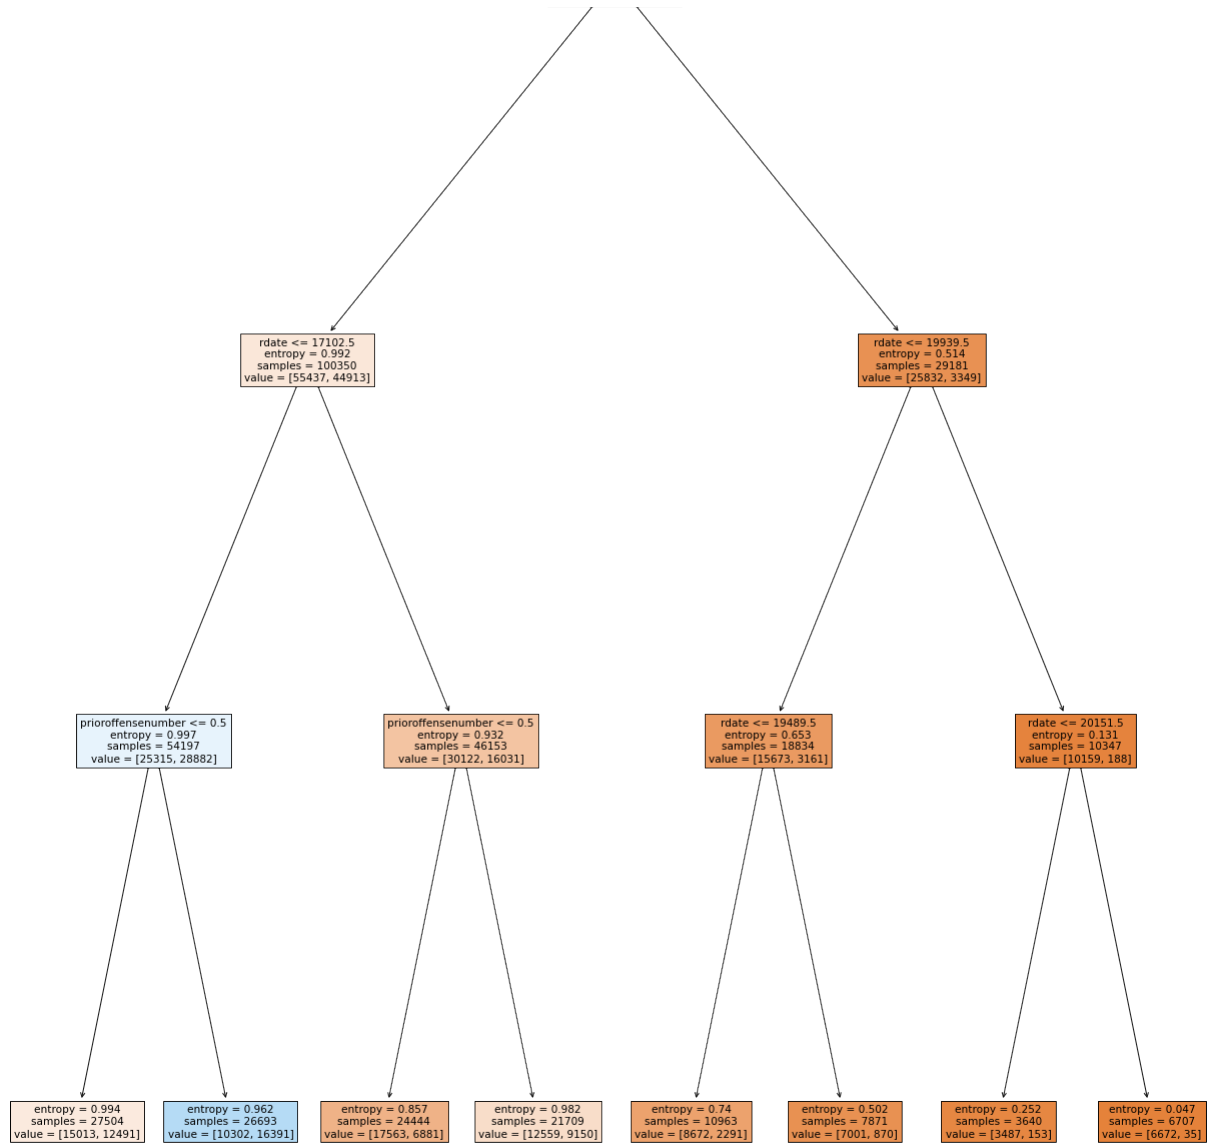
```
In [27]: # Accuracy
print('Accuracy score:', accuracy_score(Y, clf.predict(X)))

# Residual value
resid_dev = np.sum(log_loss(Y, clf.predict_proba(X)))
print('Residual value:', resid_dev)
```

Accuracy score: 0.6744177069581799
Residual value: 0.5810712818689788

```
In [28]: # Plot the tree
feature_names = X.columns
ax = subplots(figsize=(22,28))[1]
plot_tree(clf,
          feature_names=feature_names,
          ax=ax, filled = True);
```

```
maxdate <= 18907.5
entropy = 0.953
samples = 129531
value = [81269, 48262]
```



```

In [29]: # Convert feature_names (pandas Index) to a list
feature_names = feature_names.tolist()

# Print the decision tree rules
tree_rules = export_text(clf, feature_names=feature_names, show_weight
print(tree_rules)

|--- maxrdate <= 18907.50
|   |--- rdate <= 17102.50
|       |--- prioroffensenumbe <= 0.50
|           |--- weights: [15013.00, 12491.00] class: 0.0
|           |--- prioroffensenumbe > 0.50
|               |--- weights: [10302.00, 16391.00] class: 1.0
|       |--- rdate > 17102.50
|           |--- prioroffensenumbe <= 0.50
|               |--- weights: [17563.00, 6881.00] class: 0.0
|               |--- prioroffensenumbe > 0.50
|                   |--- weights: [12559.00, 9150.00] class: 0.0
|   |--- maxrdate > 18907.50
|       |--- rdate <= 19939.50
|           |--- rdate <= 19489.50
|               |--- weights: [8672.00, 2291.00] class: 0.0
|               |--- rdate > 19489.50
|                   |--- weights: [7001.00, 870.00] class: 0.0
|       |--- rdate > 19939.50
|           |--- rdate <= 20151.50
|               |--- weights: [3487.00, 153.00] class: 0.0
|               |--- rdate > 20151.50
|                   |--- weights: [6672.00, 35.00] class: 0.0

```

```

In [30]: # Pruning Classification Tree
# Load the packages
import sklearn.model_selection as skm
import matplotlib.pyplot as plt
from matplotlib.pyplot import subplots
from sklearn.tree import (DecisionTreeClassifier as DTC,
                           DecisionTreeRegressor as DTR,
                           plot_tree,
                           export_text)
from sklearn.metrics import (accuracy_score,
                              log_loss)
from sklearn.ensemble import \
    (RandomForestRegressor as RF,
     GradientBoostingRegressor as GBR,
     RandomForestClassifier as RFC)

```

```
In [31]: # Split the data between train and text
(X_train,
 X_test,
 Y_train,
 Y_test) = skm.train_test_split(X,
                                Y,
                                test_size=0.3,
                                random_state=0)

# Accuracy score
clf1 = DTC(criterion='entropy', random_state=0)
clf1.fit(X_train, Y_train)
accuracy_score(Y_test, clf.predict(X_test))
```

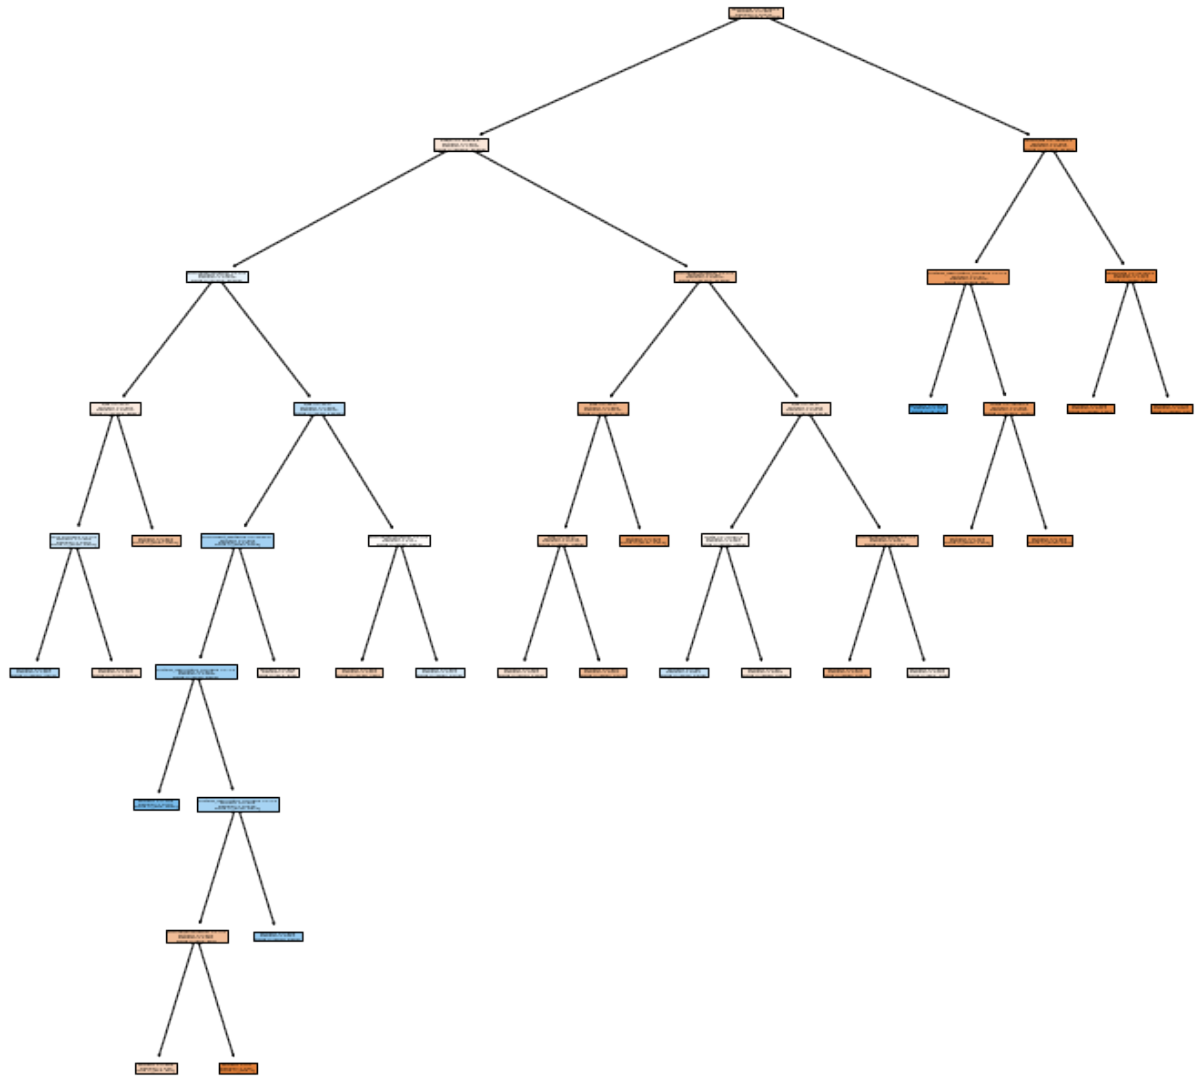
Out[31]: 0.674704065877509

```
In [32]: ccp_path = clf.cost_complexity_pruning_path(X_train, Y_train)
kfold1 = skm.KFold(10,
                  random_state=1,
                  shuffle=True)

grid = skm.GridSearchCV(clf1,
                        {'ccp_alpha': ccp_path.ccp_alphas},
                        refit=True,
                        cv=kfold1,
                        scoring='accuracy')
grid.fit(X_train, Y_train)
grid.best_score_
```

Out[32]: 0.6982829875656346

```
In [33]: # Plot the pruned tree
ax = subplots(figsize=(12, 12))[1]
best_ = grid.best_estimator_
plot_tree(best_,
          feature_names=feature_names,
          ax=ax, filled=True);
```




```
In [35]: # Building the forest
data_RF = RFC(max_features = X_train.shape[1], random_state=0)
data_RF.fit(X_train,Y_train)
```

```
Out[35]: RandomForestClassifier
RandomForestClassifier(max_features=84, random_state=0)
```

```
In [36]: # Table importance
feature_imp = pd.DataFrame(
    {'importance':data_RF.feature_importances_},
    index=feature_names)
feature_imp.sort_values(by='importance', ascending=False)
```

```
Out[36]:
```


	importance
rdate	0.086171
maxrdate	0.081360
age	0.065387
lincome	0.050155
maxdate	0.047501
...	...
traffher	0.000000
traffamph	0.000000
traffill	0.000000
traffconspir	0.000000
traffcoc	0.000000

84 rows × 1 columns

```
In [37]: # Plot variable importance
importance = data_RF.feature_importances_
sorted_importance = importance.argsort()

plt.figure(figsize=(10, 16))
plt.barh(X.columns[sorted_importance], importance[sorted_importance],c
plt.xlabel('Random Forest Feature Importance')
plt.title('Feature Importance Plot')
plt.show()
```

Feature Importance Plot



```
In [38]: import numpy as np
import matplotlib.pyplot as plt
import pandas as pd

# Assuming data_RF is your fitted Random Forest model and X is your fe
importance2 = data_RF.feature_importances_
sorted_importance2 = np.argsort(importance2)
```

```

# Create a DataFrame for better handling
feature_importance_df2 = pd.DataFrame({
    'Feature': X.columns[sorted_importance2],
    'Importance': importance[sorted_importance2]
})
# Dictionary for renaming multiple variables
rename_dict = {
    'unemp_rate': 'unemployment',
    'totalyearssentenced': 'sentence length',
    'fincrime': 'financial crime',
    'priooffensenum': 'number of prior offenses',
    'birthyear': 'birth year',
    'county1_encoded': 'county',
    'releasemonth': 'month of release',
    'rdate': 'Date released from prison',
    'lncome': 'log of income',
    'percentage_snap_recipients': 'probability of receiving SNAP'
}
# Change 'unemp_rate' to 'employment'
feature_importance_df2['Feature'] = feature_importance_df2['Feature'].replace(rename_dict)


# Highlighted features
highlighted_features = ['age', 'unemployment', 'male', 'log of income', 'probability of receiving SNAP']


# Create a color array: Use a bright color for highlighted features and grey for others
colors = ['red' if feature in highlighted_features else 'grey' for feature in feature_importance_df2['Feature']]

plt.figure(figsize=(10, 16))
plt.barh(feature_importance_df2['Feature'], feature_importance_df2['Importance'], color=colors)
plt.xlabel('Random Forest Feature Importance', fontsize=14)
plt.title('Feature Importance Plot', fontsize=16)
plt.grid(axis='x', linestyle='--', alpha=0.7)

# Show plot
plt.savefig('importance.png', format='png', dpi=300, bbox_inches='tight')
plt.show();

```





```
In [39]: ##### SUMMARY STATISTICS #####
```

```
In [40]: # Adjusting column names to match the dataset and re-running the analysis
columns_of_interest = [ 'income', 'percentage_snap_recipients', 'anyrecid' ]

# Filtering the dataset for the relevant columns and grouping by 'anyrecid'
summary_stats = data[columns_of_interest].groupby('anyrecid').describe()

# Display the summary statistics
print(summary_stats)
```

```

              income
\
      count      mean      std      min      25%      50%
anyrecid
0.0      81269.0  43970.506048  5890.209282  24031.0  40589.0  44118.0
1.0      48262.0  42110.100141  5792.292067  23852.0  38457.0  42311.0

              percentage_snap_recipients
\
      75%      max      count      mean
anyrecid
0.0      47876.0  67967.0      81269.0  0.262735  0.118334
1.0      45995.0  61379.0      48262.0  0.185716  0.089515

              min      25%      50%      75%      max
anyrecid
0.0      0.109079  0.142761  0.246107  0.393976  0.417409
1.0      0.109079  0.131741  0.147436  0.187432  0.417409
```

```
In [41]: ##### PAPER REGRESSIONS #####
```

```
In [42]: # Load needed packages
import statsmodels.formula.api as smf
from sklearn.linear_model import LogisticRegression
```

```
In [43]: # Regression paper
reg = smf.ols('Y ~ after + dist + after * dist + unemp_rate + after *
              print(reg.summary())
```

OLS Regression Results

```
=====
=====
```

Dep. Variable:	Y	R-squared:			
0.114					
Model:	OLS	Adj. R-squared:			
0.114					
Method:	Least Squares	F-statistic:			
2373.					
Date:	Fri, 20 Dec 2024	Prob (F-statistic):			
0.00					
Time:	20:13:10	Log-Likelihood:			
-81850.					
No. Observations:	129531	AIC:			
1.637e+05					
Df Residuals:	129523	BIC:			
1.638e+05					
Df Model:	7				
Covariance Type:	nonrobust				
=====					
=====					
		coef	std err	t	P> t
	[0.025	0.975]			

Intercept		1.6524	0.111	14.874	0.00
0	1.435	1.870			
after		-0.0944	0.025	-3.780	0.00
0	-0.143	-0.045			
dist		-4.374e-05	7.71e-05	-0.567	0.57
1	-0.000	0.000			
after:dist		2.109e-05	7.72e-05	0.273	0.78
5	-0.000	0.000			
unemp_rate		-0.0085	0.005	-1.718	0.08
6	-0.018	0.001			
after:unemp_rate		0.0226	0.005	4.540	0.00
0	0.013	0.032			
lincome		-0.0841	0.010	-8.203	0.00
0	-0.104	-0.064			
percentage_snap_recipients		-1.3006	0.024	-53.121	0.00
0	-1.349	-1.253			
=====					
=====					
Omnibus:	977370.264	Durbin-Watson:			
1.875					
Prob(Omnibus):	0.000	Jarque-Bera (JB):			
13963.171					
Skew:	0.363	Prob(JB):			
0.00					
Kurtosis:	1.565	Cond. No.			
4.27e+05					
=====					
=====					

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, $4.27e+05$. This might indicate that there are strong multicollinearity or other numerical problems.

In [44]: ##### DIRECTED ACYCLIC GRAPH #####

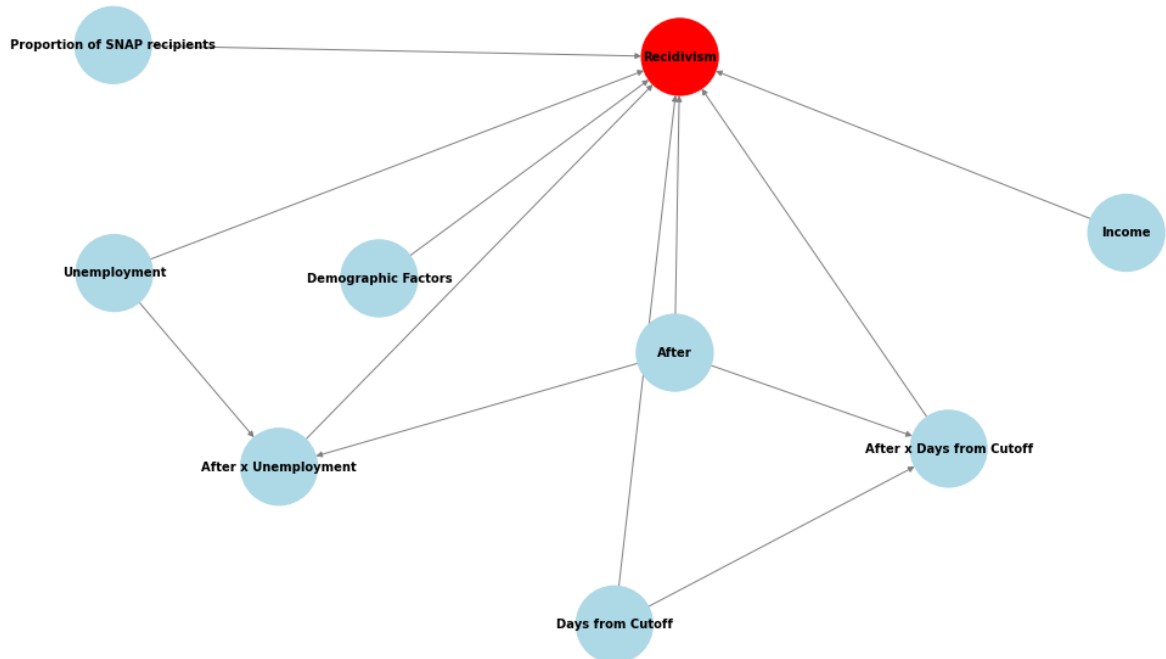
```
In [49]: # Importing necessary libraries
import matplotlib.pyplot as plt
import networkx as nx

# Create a directed acyclic graph (DAG)
G = nx.DiGraph()

# Adding nodes for variables
nodes = [
    "Recidivism", "After", "Days from Cutoff", "Unemployment", "Income",
    "Proportion of SNAP recipients", "After x Days from Cutoff", "After x Unemployment",
    "Demographic Factors"
]
G.add_nodes_from(nodes)

# Adding edges to represent relationships based on the equation
edges = [
    ("After", "Recidivism"),
    ("Days from Cutoff", "Recidivism"),
    ("Unemployment", "Recidivism"),
    ("Income", "Recidivism"),
    ("Proportion of SNAP recipients", "Recidivism"),
    ("After", "After x Days from Cutoff"),
    ("Days from Cutoff", "After x Days from Cutoff"),
    ("After x Days from Cutoff", "Recidivism"),
    ("After", "After x Unemployment"),
    ("Unemployment", "After x Unemployment"),
    ("After x Unemployment", "Recidivism"),
    ("Demographic Factors", "Recidivism")
]
G.add_edges_from(edges)

# Plotting the DAG
plt.figure(figsize=(14, 8))
pos = nx.spring_layout(G, seed=42)
node_colors = ["red" if node == "Recidivism" else "lightblue" for node in nodes]
nx.draw(G, pos, with_labels=True, node_size=4000, node_color=node_colors,
        edge_color="black", arrows=True, width=2)
plt.savefig('Directed_DAG.png', format='png', dpi=300, bbox_inches='tight')
plt.show()
```



```

In [50]: # Importing necessary libraries
import matplotlib.pyplot as plt
import networkx as nx

# Create a directed acyclic graph (DAG)
G1 = nx.DiGraph()

# Adding nodes for variables based on dataset column names
nodes = [
    "anyrecid", "after", "dist", "unemp_rate", "income",
    "percentage_snap_recipients", "after x dist", "after x unemp_rate",
    "Demographic Factors (e.g., race, sex, age)"
]
G1.add_nodes_from(nodes)

# Adding edges to represent relationships based on the equation
edges = [
    ("after", "anyrecid"),
    ("dist", "anyrecid"),
    ("unemp_rate", "anyrecid"),
    ("income", "anyrecid"),
    ("percentage_snap_recipients", "anyrecid"),
    ("after", "after x dist"),
    ("dist", "after x dist"),
    ("after x dist", "anyrecid"),
    ("after", "after x unemp_rate"),
    ("unemp_rate", "after x unemp_rate"),
  ]

```



```

    ("after x unemp_rate", "anyrecid"),
    ("Demographic Factors (e.g., race, sex, age)", "anyrecid")
]
G1.add_edges_from(edges)

# Plotting the DAG
plt.figure(figsize=(14, 8))
pos = nx.spring_layout(G1, seed=42)
nx.draw(G1, pos, with_labels=True, node_size=3000, node_color="lightblue",
        plt.title("Directed Acyclic Graph (DAG) for Recidivism Model (Updated
plt.show()

```

Directed Acyclic Graph (DAG) for Recidivism Model (Updated with Dataset Variables)



In [51]: ##### BACKDOOR #####

In [52]: # Load packages
from dowhy import CausalModel

```
In [53]: # Setting causal model
model = CausalModel(data=data,
                    treatment='after',
                    outcome='anyrecid',
                    common_causes=['unemp_rate', 'lincome', 'percentage_
graph=G1)

# Identify the estimand
estimand = model.identify_effect()
```

/Users/edithsimochemo/opt/anaconda3/lib/python3.9/site-packages/dowh
y/causal_model.py:582: UserWarning: 3 variables are assumed unobserve
d because they are not in the dataset. Configure the logging level to
`logging.WARNING` or higher for additional details.
warnings.warn(

```
In [54]: import warnings

# Display the identified estimand
print(estimand)

# Suppress all warnings
with warnings.catch_warnings():
    warnings.simplefilter("ignore")

    # Obtain estimates
    estimate = model.estimate_effect(identified_estimand=estimand,
                                    method_name='backdoor.linear_regr
    print(estimate)
```

Estimand type: EstimandType.NONPARAMETRIC_ATE

Estimand : 1

Estimand name: backdoor

Estimand expression:

$$\frac{d}{E[\text{anyrecid}]}$$

$d[\text{after}]$

Estimand assumption 1, Unconfoundedness: If $U \rightarrow \{\text{after}\}$ and $U \rightarrow \text{anyrecid}$
then $P(\text{anyrecid}|\text{after}, U) = P(\text{anyrecid}|\text{after})$

Estimand : 2

Estimand name: iv

No such variable(s) found!

Estimand : 3

Estimand name: frontdoor

No such variable(s) found!

*** Causal Estimate ***

Identified estimand

Estimand type: EstimandType.NONPARAMETRIC_ATE

Estimand : 1

Estimand name: backdoor

Estimand expression:

$$\frac{d}{d[after]}(E[anyrecid])$$

Estimand assumption 1, Unconfoundedness: If $U \rightarrow \{after\}$ and $U \rightarrow anyrecid$ then $P(anyrecid|after,,U) = P(anyrecid|after,)$

Realized estimand

b: anyrecid~after+after*unemp_rate+after*dist+after*percentage_snap_recipients+after*income

Target units:

Estimate

Mean value: -0.18008907485851255

Conditional Estimates

__categorical__unemp_rate	__categorical__dist	__categorical__percentage_snap_recipients	__categorical__income
(1.399, 3.3]	(-0.001, 1205.0]	(0.108, 0.135]	

(23851.999, 38477.0]	-0.003388
----------------------	-----------

(38477.0, 41960.0]	-0.008728
--------------------	-----------

(41960.0, 44850.0]	-0.011889
--------------------	-----------

(44850.0, 48478.0]	-0.023568
--------------------	-----------

(48478.0, 67967.0]	-0.026729
--------------------	-----------

(8.9, 22.6]	(4576.0, 6857.0]	(0.389, 0.417]
-------------	------------------	----------------

(23851.999, 38477.0]	-0.366911
----------------------	-----------

(38477.0, 41960.0]	-0.386116
--------------------	-----------

(41960.0, 44850.0]	-0.390642
--------------------	-----------

(44850.0, 48478.0]	-0.396940
--------------------	-----------

(48478.0, 67967.0]	-0.410806
--------------------	-----------

Length: 391, dtype: float64

In [55]: ##### INVERSE PROBABILITY WEIGHTING #####

```

In [69]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler

df = pd.read_stata('final_fl2.dta')

def calculate_propensity_scores(df, treatment, covariates, outcome):
    """
    Calculate propensity scores and weights for causal inference analysis.

    Parameters:
    -----
    df : pandas.DataFrame
        Input dataset
    treatment : str
        Name of treatment variable column
    covariates : list
        List of covariate column names
    outcome : str
        Name of outcome variable column

    Returns:
    -----
    tuple
        (DataFrame with weights, propensity scores array, ATE estimate)
    """
    # Create a copy to avoid modifying original data
    df_clean = df.copy()

    # Drop rows with missing values
    df_clean = df_clean.dropna(subset=[treatment] + covariates + [outcome])

    # Standardize covariates
    scaler = StandardScaler()
    X = scaler.fit_transform(df_clean[covariates])
    X = pd.DataFrame(X, columns=covariates, index=df_clean.index)
    y = df_clean[treatment]

    # Fit logistic regression with balanced class weights
    logit = LogisticRegression(class_weight='balanced', random_state=42)
    logit.fit(X, y)

    # Calculate propensity scores
    propensity_scores = logit.predict_proba(X)[:, 1]

    # Trim extreme propensity scores to avoid infinite weights

```

```

eps = 0.01
propensity_scores = np.clip(propensity_scores, eps, 1 - eps)

# Calculate inverse probability weights
weights = np.where(y == 1,
                    1/propensity_scores,
                    1/(1 - propensity_scores))

# Add scores and weights to dataframe
df_clean['propensity_score'] = propensity_scores
df_clean['weights'] = weights

# Calculate ATE
ate = np.average(df_clean[outcome], weights=df_clean['weights'])

return df_clean, propensity_scores, ate

def plot_positivity_check(df, treatment_col='after'):
    """
    Create positivity check plot for propensity scores.

    Parameters:
    -----
    df : pandas.DataFrame
        DataFrame containing propensity scores
    treatment_col : str
        Name of treatment variable column
    """
    plt.figure(figsize=(10, 6))

    # Create separate density plots for each group
    for group in [0, 1]:
        mask = df[treatment_col] == group
        plt.hist(df.loc[mask, 'propensity_score'],
                 bins=30,
                 density=True,
                 alpha=0.5,
                 label=f"{'Treatment' if group == 1 else 'Control'}")

    plt.title('Positivity Check: Distribution of Propensity Scores by')
    plt.xlabel('Propensity Score')
    plt.ylabel('Density')
    plt.legend()
    plt.grid(True, alpha=0.3)
    plt.savefig('positivity_check.png', format = 'png', dpi =300, bbox
    plt.show()

if __name__ == "__main__":
    # Define variables

```

```

treatment = 'after'
covariates = ['age', 'male', 'black',
              'prioroffense', 'unemp_rate', 'lincome', 'percentage_
outcome = 'anyrecid'

# Calculate propensity scores and weights
df_weighted, prop_scores, ate = calculate_propensity_scores(
    df, treatment, covariates, outcome
)

# Display results
print("\nSample of propensity scores and weights:")
print(df_weighted[['offenderid', treatment, 'propensity_score', 'w

print(f"\nEstimated Average Treatment Effect (ATE): {ate:.4f}")

# Create positivity check plot
plot_positivity_check(df_weighted, treatment)

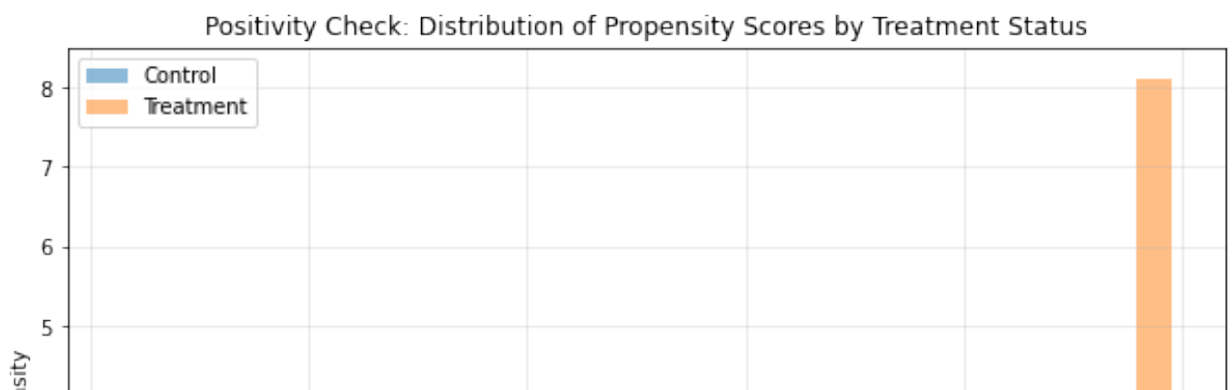
# Print covariate balance summary
print("\nCovariate balance summary:")
for cov in covariates:
    treated_mean = np.average(df_weighted.loc[df_weighted[treatment
                                weights=df_weighted.loc[df_weighted[tr
    control_mean = np.average(df_weighted.loc[df_weighted[treatment
                                weights=df_weighted.loc[df_weighted[tr
    print(f"{cov}: Treated mean = {treated_mean:.2f}, Control mean


```

Sample of propensity scores and weights:

	offenderid	after	propensity_score	weights
1	A000043	1.0	0.137165	7.290506
2	A000043	1.0	0.938853	1.065130
3	A000077	1.0	0.368211	2.715831
5	A000093	0.0	0.393505	1.648819
6	A000101	1.0	0.516569	1.935851

Estimated Average Treatment Effect (ATE): 0.3769





Covariate balance summary:

age: Treated mean = 32.26, Control mean = 33.30

male: Treated mean = 0.90, Control mean = 0.92

black: Treated mean = 0.52, Control mean = 0.50

prioroffense: Treated mean = 0.38, Control mean = 0.37

unemp_rate: Treated mean = 5.27, Control mean = 5.68

lincome: Treated mean = 10.59, Control mean = 10.64

percentage_snap_recipients: Treated mean = 0.20, Control mean = 0.24

In [63]: ##### METALEARNERS #####

```
In [74]: # Import packages
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

df = pd.read_stata('final_fl2.dta')

# Modify features based on available columns
features = ['concurrent_sentence', 'drugoffense', 'traffoffense', 'oth
            'age', 'male', 'black', 'totalyears sentenced', 'prioroffens
            'lincome', 'percentage_snap_recipients']
treatment = 'after'
outcome = 'anyrecid'

# Drop rows with missing values
df_clean = data[features + [treatment, outcome]].dropna()

# Split features
X = df_clean[features]
T = df_clean[treatment]
Y = df_clean[outcome]

# S-Learner
```

```
def s_learner(X, T, Y):
    # Combine features with treatment
    X_s = X.copy()
    X_s['treatment'] = T

    # Train random forest
    s_model = RandomForestRegressor(n_estimators=100, random_state=42)
    s_model.fit(X_s, Y)

    # Predict potential outcomes
    X_s_1 = X_s.copy()
    X_s_0 = X_s.copy()
    X_s_1['treatment'] = 1
    X_s_0['treatment'] = 0

    y1_pred = s_model.predict(X_s_1)
    y0_pred = s_model.predict(X_s_0)

    # Calculate ATE
    ate = np.mean(y1_pred - y0_pred)
    return ate, y1_pred, y0_pred

# T-Learner
def t_learner(X, T, Y):
    # Split data by treatment
    X_t = X[T == 1]
    X_c = X[T == 0]
    Y_t = Y[T == 1]
    Y_c = Y[T == 0]

    # Train separate models
    t1_model = RandomForestRegressor(n_estimators=100, random_state=42)
    t0_model = RandomForestRegressor(n_estimators=100, random_state=42)

    t1_model.fit(X_t, Y_t)
    t0_model.fit(X_c, Y_c)

    # Predict potential outcomes
    y1_pred = t1_model.predict(X)
    y0_pred = t0_model.predict(X)

    # Calculate ATE
    ate = np.mean(y1_pred - y0_pred)
    return ate, y1_pred, y0_pred

# X-Learner
def x_learner(X, T, Y):
    # First stage: T-Learner
    t_ate, y1_pred, y0_pred = t_learner(X, T, Y)
```



```

# Second stage: Calculate individual treatment effects
X_t = X[T == 1]
X_c = X[T == 0]
Y_t = Y[T == 1]
Y_c = Y[T == 0]

# Calculate residuals
D1 = Y_t - y0_pred[T == 1]
D0 = y1_pred[T == 0] - Y_c

# Train second stage models
x1_model = RandomForestRegressor(n_estimators=100, random_state=42)
x0_model = RandomForestRegressor(n_estimators=100, random_state=42)

x1_model.fit(X_t, D1)
x0_model.fit(X_c, D0)

# Predict treatment effects
tau1 = x1_model.predict(X)
tau0 = x0_model.predict(X)

# Calculate final treatment effect
g = np.mean(T) # Propensity score (simplified)
tau = g * tau0 + (1 - g) * tau1

return np.mean(tau), tau

# Run all models
print("Running models...")
s_ate, s_y1, s_y0 = s_learner(X, T, Y)
t_ate, t_y1, t_y0 = t_learner(X, T, Y)
x_ate, x_tau = x_learner(X, T, Y)

# Create results summary
results = pd.DataFrame({
    'Model': ['S-Learner', 'T-Learner', 'X-Learner'],
    'Average Treatment Effect': [s_ate, t_ate, np.mean(x_tau)]
})

print("\n
Model Results:")
print(results)

# Visualize treatment effects distribution
plt.figure(figsize=(10, 6))
plt.hist(s_y1 - s_y0, bins=50, alpha=0.5, label='S-Learner')
plt.hist(t_y1 - t_y0, bins=50, alpha=0.5, label='T-Learner')
plt.hist(x_tau, bins=50, alpha=0.5, label='X-Learner')
plt.xlabel('Individual Treatment Effects')
plt.ylabel('Frequency')

```

```

plt.legend()
plt.ylim(0, 30000)
plt.savefig('metalearners.png', format = 'png', dpi =300, bbox_inches=
plt.show()

# Calculate model performance metrics
def calculate_metrics(y_true, y_pred_treated, y_pred_control):
    mse_treated = mean_squared_error(y_true[T == 1], y_pred_treated[T
    mse_control = mean_squared_error(y_true[T == 0], y_pred_control[T
    return np.sqrt(mse_treated), np.sqrt(mse_control)

s_rmse_t, s_rmse_c = calculate_metrics(Y, s_y1, s_y0)
t_rmse_t, t_rmse_c = calculate_metrics(Y, t_y1, t_y0)

metrics = pd.DataFrame({
    'Model': ['S-Learner', 'T-Learner'],
    'RMSE (Treated)': [s_rmse_t, t_rmse_t],
    'RMSE (Control)': [s_rmse_c, t_rmse_c]
})

print("\
Model Performance Metrics:")
print(metrics)

# Feature importance for S-Learner
s_model = RandomForestRegressor(n_estimators=100, random_state=42)
X_s = X.copy()
X_s['treatment'] = T
s_model.fit(X_s, Y)

feature_importance = pd.DataFrame({
    'Feature': features + ['treatment'],
    'Importance': s_model.feature_importances_
}).sort_values('Importance', ascending=False)

print("\
Feature Importance (S-Learner):")
print(feature_importance)


```

Running models...

Model Results:

	Model	Average Treatment Effect
0	S-Learner	0.002248
1	T-Learner	0.000301
2	X-Learner	0.076148





Model Performance Metrics:

	Model	RMSE (Treated)	RMSE (Control)
0	S-Learner	0.173536	0.186604
1	T-Learner	0.173503	0.189568

Feature Importance (S-Learner):

	Feature	Importance
11	percentage_snap_recipients	0.258495
4	age	0.161238
10	lincome	0.160923
9	unemp_rate	0.156308
0	concurrent_sentence	0.101182
7	totalyearssentenced	0.100694
6	black	0.021428
8	prioroffense	0.021363
5	male	0.013172
12	treatment	0.005197
1	drugoffense	0.000000
2	traffoffense	0.000000
3	otheroffense	0.000000

In [73]: ##### BOOTSTRAP #####

```
In [79]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
from typing import Tuple, List, Optional
import warnings

def calculate_propensity_scores(df: pd.DataFrame,
```

```

        treatment: str,
        covariates: List[str],
        outcome: str) -> Tuple[pd.DataFrame, np.
"""
Calculate propensity scores and weights for causal inference analysis

Parameters and returns same as before
"""
# Previous implementation remains the same
df_clean = df.copy()
df_clean = df_clean.dropna(subset=[treatment] + covariates + [outcome])

scaler = StandardScaler()
X = scaler.fit_transform(df_clean[covariates])
X = pd.DataFrame(X, columns=covariates, index=df_clean.index)
y = df_clean[treatment]

logit = LogisticRegression(class_weight='balanced', random_state=42)
logit.fit(X, y)

propensity_scores = logit.predict_proba(X)[:, 1]
eps = 0.01
propensity_scores = np.clip(propensity_scores, eps, 1 - eps)

weights = np.where(y == 1,
                    1/propensity_scores,
                    1/(1 - propensity_scores))

df_clean['propensity_score'] = propensity_scores
df_clean['weights'] = weights

ate = np.average(df_clean[outcome], weights=weights)

return df_clean, propensity_scores, ate

def bootstrap_ate(df: pd.DataFrame,
                  treatment: str,
                  outcome: str,
                  covariates: List[str],
                  n_bootstrap: int = 1000,
                  random_state: Optional[int] = None) -> Tuple[float, float]:
    """
    Perform bootstrap analysis of Average Treatment Effect (ATE).

    Parameters:
    -----
    df : pandas.DataFrame
        Input dataset
    treatment : str
        Name of treatment variable column
    """

```

```

outcome : str
    Name of outcome variable column
covariates : List[str]
    List of covariate column names
n_bootstrap : int, optional
    Number of bootstrap iterations (default: 1000)
random_state : int, optional
    Random seed for reproducibility

Returns:
-----
Tuple[float, float, float, np.ndarray]
    (original ATE, lower CI, upper CI, bootstrap ATEs)
"""
if random_state is not None:
    np.random.seed(random_state)

# Calculate original ATE
df_weighted, _, original_ate = calculate_propensity_scores(
    df, treatment, covariates, outcome
)

# Perform bootstrap
bootstrap_ates = []
for _ in range(n_bootstrap):
    # Sample with replacement
    bootstrap_indices = np.random.choice(
        len(df_weighted),
        size=len(df_weighted),
        replace=True
    )
    bootstrap_sample = df_weighted.iloc[bootstrap_indices]

    # Recalculate propensity scores and weights for bootstrap sample
    _, _, bootstrap_ate = calculate_propensity_scores(
        bootstrap_sample,
        treatment,
        covariates,
        outcome
    )
    bootstrap_ates.append(bootstrap_ate)

# Calculate confidence intervals
ci_lower, ci_upper = np.percentile(bootstrap_ates, [2.5, 97.5])

return original_ate, ci_lower, ci_upper, np.array(bootstrap_ates)

def plot_bootstrap_results(original_ate: float,
                           bootstrap_ates: np.ndarray,
                           ci_lower: float,

```

```

                                ci_upper: float) -> None:
"""
Plot bootstrap distribution with confidence intervals.

Parameters:
-----
original_ate : float
    Original ATE estimate
bootstrap_ates : numpy.ndarray
    Array of bootstrap ATE estimates
ci_lower : float
    Lower bound of confidence interval
ci_upper : float
    Upper bound of confidence interval
"""
plt.figure(figsize=(10, 6))

# Plot histogram of bootstrap estimates
plt.hist(bootstrap_ates, bins=50, density=True, alpha=0.6,
        label='Bootstrap Distribution')

# Add vertical lines for original ATE and CIs
plt.axvline(original_ate, color='red', linestyle='dashed',
        label=f'Original ATE: {original_ate:.3f}')
plt.axvline(ci_lower, color='green', linestyle='dashed',
        label=f'95% CI: ({ci_lower:.3f}, {ci_upper:.3f})')
plt.axvline(ci_upper, color='green', linestyle='dashed')


plt.xlabel('ATE')
plt.ylabel('Density')
plt.legend()
plt.grid(True, alpha=0.3)
plt.savefig('bootstrap_int.png', format='png', dpi=300, bbox_inches='tight')
plt.show()

# Example usage:
if __name__ == "__main__":
    # Define variables
    treatment = 'after'
    covariates = ['age', 'male', 'black',
                  'prioroffense', 'unemp_rate', 'lincome', 'percentage_
    outcome = 'anyrecid'

    # Perform bootstrap analysis
    original_ate, ci_lower, ci_upper, bootstrap_ates = bootstrap_ate(
        df,
        treatment,
        outcome,
        covariates,

```

```
n_bootstrap=1000,  
random_state=42  
)  
  
# Plot results  
plot_bootstrap_results(original_ate, bootstrap_ates, ci_lower, ci_upper,  
                        ci_upper, ci_lower)  
  
# Print numerical results  
print(f"Original ATE: {original_ate:.3f}")  
print(f"95% Confidence Interval: ({ci_lower:.3f}, {ci_upper:.3f})")  
print(f"Standard Error: {np.std(bootstrap_ates):.3f}")  
  
# Additional statistics  
print(f"\nBootstrap Statistics:")  
print(f"Mean of bootstrap estimates: {np.mean(bootstrap_ates):.3f}")  
print(f"Median of bootstrap estimates: {np.median(bootstrap_ates):.3f}")  
print(f"Standard deviation of bootstrap estimates: {np.std(bootstrap_ates):.3f}")
```



Original ATE: 0.377
95% Confidence Interval: (0.374, 0.379)
Standard Error: 0.001

Bootstrap Statistics:
Mean of bootstrap estimates: 0.377
Median of bootstrap estimates: 0.377
Standard deviation of bootstrap estimates: 0.001

In [76]: ##### DOUBLE ML MODELS #####

```

In [85]: # librairies
from econml.dml import DML
from sklearn.linear_model import RidgeCV, LinearRegression

# Set up features and clean data
features = ['age', 'black', 'male', 'prioroffense', 'totalyearsentence',
            'prioroffensenum', 'countoffenses', 'unemp_rate', 'lincome']
df_clean = df.dropna(subset=['after', 'anyrecid'] + features)

# Prepare variables
X = df_clean[features]
T = df_clean['after'].astype(float)
Y = df_clean['anyrecid'].astype(float)

print("Data shape after cleaning:", X.shape)
print("\
Treatment variable (after) statistics:")
print(T.describe())
print("\
Outcome variable statistics:")
print(Y.describe())

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
X_scaled = pd.DataFrame(X_scaled, columns=X.columns)

# Linear Double ML Model
dml_linear = DML(
    model_y=LassoCV(random_state=123, max_iter=2000),
    model_t=LassoCV(random_state=123, max_iter=2000),
    model_final=LinearRegression(fit_intercept=True),
    random_state=123
)

# Fit and get linear model results
dml_linear.fit(Y, T, X=X_scaled)
effect_linear = dml_linear.effect(X=X_scaled)

print("\
Linear Double ML Results:")
print("Average Treatment Effect:", np.mean(effect_linear))

# Non-Linear Double ML Model
dml_nonlinear = DML(
    model_y=RandomForestRegressor(n_estimators=200, max_depth=5, random_state=123),
    model_t=RandomForestRegressor(n_estimators=200, max_depth=5, random_state=123),
    model_final=LinearRegression(fit_intercept=True),
    random_state=123
)

```



```
)

# Fit and get non-linear model results
dml_nonlinear.fit(Y, T, X=X_scaled)
effect_nonlinear = dml_nonlinear.effect(X=X_scaled)

print("\n
Non-Linear Double ML Results:")
print("Average Treatment Effect:", np.mean(effect_nonlinear))

# Compare the models
print("\n
Comparison:")
print("Difference in ATE (Non-linear - Linear):", np.mean(effect_nonli
```

Data shape after cleaning: (404821, 10)

Treatment variable (after) statistics:

count	404821.000000
mean	0.962504
std	0.189973
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	1.000000

Name: after, dtype: float64

Outcome variable statistics:

count	404821.000000
mean	0.305318
std	0.460542
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

Name: anyrecid, dtype: float64

The final model has a nonzero intercept for at least one outcome; it will be subtracted, but consider fitting a model without an intercept if possible.

Linear Double ML Results:

Average Treatment Effect: -0.08562746596567222

Non-Linear Double ML Results:

Average Treatment Effect: -0.02067909485078309

Comparison:

Difference in ATE (Non-linear - Linear): 0.06494837111488913

The final model has a nonzero intercept for at least one outcome; it will be subtracted, but consider fitting a model without an intercept if possible.

```

In [87]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

def visualize_dml_results():
    # Dataset statistics
    treatment_stats = {
        'count': 404821,
        'mean': 0.962504,
        'std': 0.189973
    }

    outcome_stats = {
        'count': 404821,
        'mean': 0.305318,
        'std': 0.460542
    }

    # Treatment effects
    effects_data = {
        'Linear DML': -0.08562746596567222,
        'Non-Linear DML': -0.02067909485078309
    }

    # Create figure with subplots
    fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))

    # Plot 1: Treatment Effects
    effects_df = pd.DataFrame(list(effects_data.items()), columns=['Model', 'Effect'])
    effects_df['Effect_Abs'] = effects_df['Effect'].abs() * 100

    sns.barplot(data=effects_df, x='Model', y='Effect_Abs', ax=ax1)
    ax1.set_title('Treatment Effects on Recidivism')
    ax1.set_ylabel('Reduction in Recidivism (%)')

    # Add value labels on bars
    for i, v in enumerate(effects_df['Effect_Abs']):
        ax1.text(i, v, f'{v:.1f}%', ha='center', va='bottom')

    # Plot 2: Sample Statistics
    stats = {
        'Treatment Group': treatment_stats['mean'] * 100,
        'Baseline Recidivism': outcome_stats['mean'] * 100
    }

    sns.barplot(x=list(stats.keys()), y=list(stats.values()), ax=ax2)
    ax2.set_title('Sample Statistics')
    ax2.set_ylabel('Percentage (%)')

```

```

# Add value labels on bars
for i, v in enumerate(stats.values()):
    ax2.text(i, v, f'{v:.1f}%', ha='center', va='bottom')

plt.tight_layout()

# Print summary statistics
print("Analysis Summary:")
print(f"Sample Size: {treatment_stats['count']:,} observations")
print(f"Treatment Group: {treatment_stats['mean']*100:.1f}% of sample")
print(f"Baseline Recidivism Rate: {outcome_stats['mean']*100:.1f}%")
print("\nTreatment Effects:")
print(f"Linear Model: {effects_data['Linear DML']*100:.2f}% reduction")
print(f"Non-Linear Model: {effects_data['Non-Linear DML']*100:.2f}% reduction")
print(f"Model Difference: {abs(effects_data['Linear DML'] - effects_data['Non-Linear DML'])*100:.2f}%")

return fig

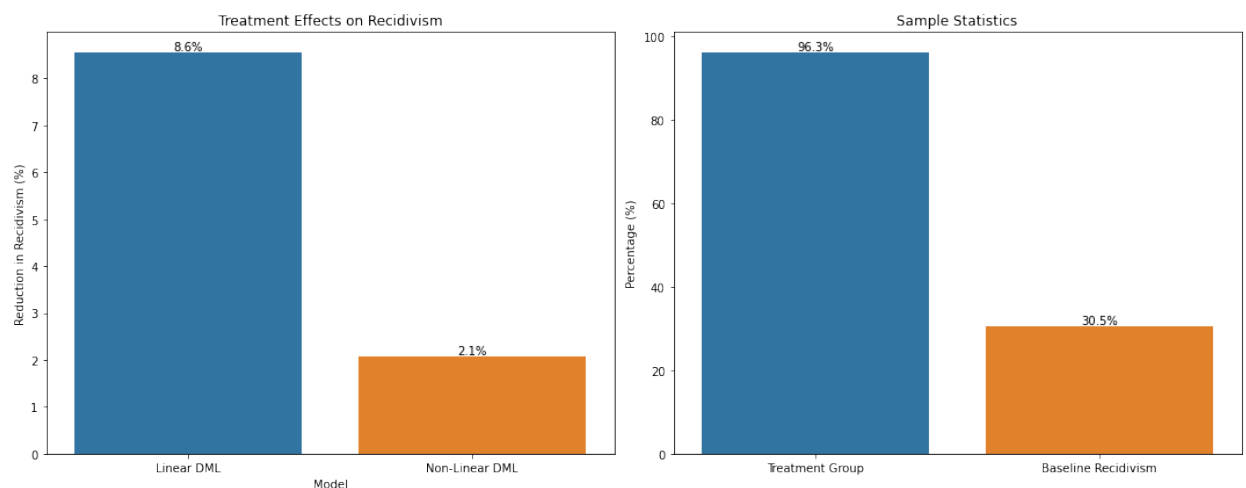
# Generate the visualization
fig = visualize_dml_results()
plt.savefig('double_robust.png', format='png', dpi=300, bbox_inches='tight')
plt.show()

```

unique with argument that is not not a Series, Index, ExtensionArray, or np.ndarray is deprecated and will raise in a future version.

Analysis Summary:
Sample Size: 404,821 observations
Treatment Group: 96.3% of sample
Baseline Recidivism Rate: 30.5%

Treatment Effects:
Linear Model: -8.56% reduction
Non-Linear Model: -2.07% reduction
Model Difference: 6.49%



```

In [94]: import numpy as np
import matplotlib.pyplot as plt

def create_clean_cumulative_gains_plot():
    # Generate sample predictions for demonstration
    np.random.seed(32)
    n_samples = 404821

    # Simulate predictions based on the given effect sizes
    linear_pred = np.random.normal(-0.08562746596567222, 0.05, n_samples)
    nonlinear_pred = np.random.normal(-0.02067909485078309, 0.05, n_samples)

    # Sort predictions in descending order
    linear_sorted = np.sort(linear_pred)[::-1]
    nonlinear_sorted = np.sort(nonlinear_pred)[::-1]

    # Calculate percentiles
    percentiles = np.arange(len(linear_sorted)) / float(len(linear_sorted))

    # Calculate cumulative gains
    linear_gains = np.cumsum(linear_sorted) / np.sum(linear_sorted)
    nonlinear_gains = np.cumsum(nonlinear_sorted) / np.sum(nonlinear_sorted)

    # Create diagonal line for random model
    diagonal = percentiles

    # Create the plot
    plt.figure(figsize=(10, 6))
    plt.plot(percentiles, linear_gains, 'b-', label='Linear DML (-8.56%)')
    plt.plot(percentiles, nonlinear_gains, 'r-', label='Non-Linear DML (-10.56%)')
    plt.plot(percentiles, diagonal, 'k--', label='Random Model', linewidth=2)

    plt.xlabel('Percentage of Sample')
    plt.ylabel('Cumulative Gain')
    #plt.title('Cumulative Gains Plot: Linear vs Non-Linear DML Models')
    plt.legend(loc='lower right')
    plt.grid(False)


    # Add text box with key statistics
    plt.text(0.05, 0.95,
            'Sample Size: 405,006\nBaseline Recidivism: 30.5%',
            transform=plt.gca().transAxes,
            bbox=dict(facecolor='white', alpha=0.8))

    return plt.gcf()

# Generate and save the plot
fig = create_clean_cumulative_gains_plot()
plt.savefig('cumulative_gain.png', format='png', dpi=300, bbox_inches='tight')

```

```
plt.show()
```



```
In [77]: ##### HETEROGENOUS TREATMENT EFFECTS #####
```

```
In [99]: import numpy as np
from econml.grf import CausalForest

# Ensure X is a 2D array
X = data.drop(columns=["anyrecid", "distn3", "distn4", "finrecidany",

# Ensure T and Y are 2D arrays with a single column
T = data["after"].values.reshape(-1, 1)
Y = data["anyrecid"].values.reshape(-1, 1)

# Fit a Generalized Random Forest
grf = CausalForest(n_estimators=500, min_samples_leaf=10, max_depth=No
grf.fit(X, T, Y)

# Estimate Conditional Average Treatment Effects (CATE)
tau_hat = grf.predict(X)
```

```
In [100]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns


# Flatten tau_hat if it's a 2D array
tau_hat_flat = tau_hat.flatten()
```

```

# Plot the histogram of heterogeneous treatment effects
plt.figure(figsize=(10, 6))
plt.hist(tau_hat_flat, bins=30, color="blue", alpha=0.7, edgecolor='black')
#plt.title("Distribution of Heterogeneous Treatment Effects (HTEs)")
plt.xlabel("HTE Estimate")
plt.ylabel("Frequency")
plt.axvline(x=np.mean(tau_hat_flat), color='red', linestyle='dashed',
            label='Mean HTE')
plt.legend()
plt.savefig('heterogeneous.png', format = 'png', dpi = 300, bbox_inches='tight')
plt.show()

# Descriptive statistics
print("Mean HTE:", np.mean(tau_hat_flat))
print("Standard Deviation of HTE:", np.std(tau_hat_flat))
print("Median HTE:", np.median(tau_hat_flat))
print("Minimum HTE:", np.min(tau_hat_flat))
print("Maximum HTE:", np.max(tau_hat_flat))

```



```

Mean HTE: 0.028699354205585472
Standard Deviation of HTE: 0.07566421105812537
Median HTE: 0.03568033397496606
Minimum HTE: -0.47547132692639593
Maximum HTE: 0.2747431966109275

```

```

In [113]: dataset = pd.read_stata("final_fl2.dta")

# Convert Categorical variables: List of categorical columns
categorical_cols = ['race', 'sex', 'custody_description', 'county1']

```

```

# Create an instance of LabelEncoder
le = LabelEncoder()

# Apply LabelEncoder to each categorical column
for col in categorical_cols:
    dataset[col+ '_encoded'] = le.fit_transform(data[col])

# Remove missing observations
df = dataset.dropna()

```

```

-----
-----
KeyError                                Traceback (most recent call
last)
~/opt/anaconda3/lib/python3.9/site-packages/pandas/core/indexes/base.
py in get_loc(self, key)
    3804         try:
-> 3805             return self._engine.get_loc(casted_key)
    3806         except KeyError as err:

index.pyx in pandas._libs.index.IndexEngine.get_loc()

index.pyx in pandas._libs.index.IndexEngine.get_loc()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.PyO
bjectHashTable.get_item()

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.PyO
bjectHashTable.get_item()

KeyError: 'race'

```

The above exception was the direct cause of the following exception:

```

KeyError                                Traceback (most recent call
last)
/var/folders/1x/fvyxnz7d5db2mw8jz8k69snr0000gn/T/ipykernel_48306/1478
161101.py in <module>
     9 # Apply LabelEncoder to each categorical column
    10 for col in categorical_cols:
--> 11     dataset[col+ '_encoded'] = le.fit_transform(data[col])
    12
    13 # Remove missing observations

~/opt/anaconda3/lib/python3.9/site-packages/pandas/core/frame.py in _
getitem__(self, key)
    4100         if self.columns.nlevels > 1:
    4101             return self._getitem_multilevel(key)

```

```

-> 4102             indexer = self.columns.get_loc(key)
    4103             if is_integer(indexer):
    4104                 indexer = [indexer]

~/opt/anaconda3/lib/python3.9/site-packages/pandas/core/indexes/base.
py in get_loc(self, key)
    3810         ):
    3811             raise InvalidIndexError(key)
-> 3812         raise KeyError(key) from err
    3813     except TypeError:
    3814         # If we have a listlike key, _check_indexing_erro
r will raise

KeyError: 'race'

```

```

In [115]: import pandas as pd
from sklearn.preprocessing import LabelEncoder

# Load dataset
dataset = pd.read_stata("final_fl2.dta")

# Print the column names to verify
print("Column Names in Dataset:", dataset.columns)

# Define categorical columns
categorical_cols = ['race', 'sex', 'custody_description', 'county1']

# Ensure column names are correct
corrected_cols = [col.strip().lower() for col in dataset.columns]
categorical_cols = [col for col in categorical_cols if col in corrected_cols]

# Encode categorical variables
le = LabelEncoder()

for col in categorical_cols:
    dataset[col + '_encoded'] = le.fit_transform(dataset[col])

# Drop missing observations
df = dataset.dropna()

# Output final dataframe
df.head()

```

```

Column Names in Dataset: Index(['level_0', 'index', 'offenderid', 'date', 'adate', 'redate', 'rdate', 'county1', 'releaseyear', 'releasemonth', 'after', 'dist', 'distnoab', 'distn2', 'distn3', 'distn4', 'fullbanafter', 'fullbanbefore', 'concurrent_sentence', 'drugoffense', 'traffoffense', 'otherof

```



```
fense',
    'smd', 'traffmar', 'traffcoc', 'traffher', 'traffamph', 'traff
ill',
    'traffconspir', 'fincrime', 'notpossoffense', 'drugoffense_nos
elling',
    'drugoffense_poss', 'violentcrime', 'assault', 'elderly', 'esc
ape',
    'forgery', 'fraud', 'kidnap', 'manslaughter', 'murder', 'other
crime',
    'otherviolent', 'propdamage', 'racketeer', 'robbery', 'sexcrim
e',
    'propsteal', 'weapon', 'criminalmischief', 'dui', 'licrevoke',
    'fleeoescape', 'fraudforge', 'anytheft', 'anyburg', 'propcrim
e',
    'race', 'sex', 'birthyear', 'custody_description',
    'facility_description', '_mergedemo', 'maxdate', 'maxadate', '
maxrdate',
    'dateorig', 'offenseyear', 'offensemonth', 'ban', 'age', 'unde
r30',
    'black', 'male', 'totalyearssentenced', 'prioroffensenumbe
r',
    'prioroffense', 'countoffenses', 'anyrecid', 'finrecidany',
    'nonfinrecidany', 'preoct97', 'placebodrug', 'placebosmd', 'ye
ar_x',
    'month_x', 'county_x', 'unemp_rate', 'year_y', 'county_y', 'in
come',
    'lincome', 'year', 'month_y', 'percentage_snap_recipients'],
dtype='object')
```

Out[115]:

	level_0	index	offenderid	date	adate	redate	rdate	county1	releaseyear	relea
3	3	3	A000077	13613.0	14005.0	1998-07-02	14661.0	PALM BEACH	2000.0	
7	7	7	A000102	17364.0	17639.0	2008-05-01	17902.0	BAY	2009.0	
19	19	19	A000470	13633.0	14202.0	1998-12-10	14838.0	SARASOTA	2000.0	
20	20	20	A000470	14888.0	15011.0	2001-02-22	15500.0	PINELLAS	2002.0	
22	22	22	A000495	15057.0	15264.0	2001-11-14	15870.0	POLK	2003.0	

5 rows × 100 columns